

HW2 - Theory

hhp9256 - Hoang Pham

1 Theory

1.1 Convolutional Neural Networks

- (a) The output size is 9×13 .
- (b) F filters will yield F channels in the output. Therefore, the output tensor size is

$$F \times \left[\frac{(H + 2P) - (H_K + D(H_K - 1))}{S} + 1 \right] \times \left[\frac{(W + 2P) - (W_K + D(W_K - 1))}{S} + 1 \right]$$

- (c) (i) Note that the convention used in PyTorch for ConvNet weight shape is always (out_channels, in_channels, kernel size), so $f_{\mathbf{W}}(\mathbf{x}) \in \mathbb{R}^5$, where $5 = \frac{11-3}{2} + 1$, and $\mathbf{W} \in \mathbb{R}^{1 \times 7 \times 3}$ with

$$f_{\mathbf{W}}(\mathbf{x})[n] = \sum_{k=1}^7 \sum_{i=1}^3 \mathbf{x}[2(n-1) + i][k] \mathbf{W}[1, k, i]$$

- (ii) Since $f_{\mathbf{W}}(\mathbf{x}) \in \mathbb{R}^5$, and $\mathbf{W} \in \mathbb{R}^{1 \times 7 \times 3}$, we have $\frac{\partial f_{\mathbf{W}}(\mathbf{x})}{\partial \mathbf{W}} \in \mathbb{R}^{5 \times (1 \times 7 \times 3)}$, and

$$\frac{\partial f_{\mathbf{W}}(\mathbf{x})}{\partial \mathbf{W}}[n, 1, k, i] = \mathbf{x}[2(n-1) + i][k].$$

- (iii) Since $\mathbf{x} \in \mathbb{R}^{7 \times 11}$, we have $\frac{\partial f_{\mathbf{W}}(\mathbf{x})}{\partial \mathbf{x}} \in \mathbb{R}^{5 \times 7 \times 11}$, and, using 1-based indexing:

$$\frac{\partial f_{\mathbf{W}}(\mathbf{x})}{\partial \mathbf{x}}[n, k, i] = \begin{cases} \mathbf{W}[1, k, i - 2(n-1)] & \text{when } 2(n-1) < i \leq 3 \\ 0 & \text{otherwise} \end{cases}$$

- (iv) Since $\mathbf{W} \in \mathbb{R}^{1 \times 7 \times 3}$, we have $\frac{\partial \ell}{\partial \mathbf{W}} \in \mathbb{R}^{1 \times 7 \times 3}$, since ℓ is a scalar. Additionally, by chain rule $\frac{\partial \ell}{\partial \mathbf{W}} = \frac{\partial \ell}{\partial f_{\mathbf{W}}(\mathbf{x})} \frac{\partial f_{\mathbf{W}}(\mathbf{x})}{\partial \mathbf{W}}$, we have:


$$\frac{\partial \ell}{\partial \mathbf{W}}[1, k, i] = \sum_{n=1}^2 \left(\frac{\partial \ell}{\partial f_{\mathbf{W}}(\mathbf{x})} \right) [n]. \mathbf{x}[2(n-1) + i][k]$$

This formula is calculating gradient for a backward pass, while the first part (i) is a forward pass. The initial forward pass performs a stride (size 2), which makes the backward pass perform a dilation operation.

1.2 Recurrent Neural Networks

1.2.1 Part 1

- (a) Attached is the graph of the RNN.



(b) The dimension of c_t is \mathbb{R}^m .

(c) For two vectors x, y , $x \odot y = D(y)x = D(x)y$, where $D(c)$ denotes diagonal matrix created from c . Therefore

$$h_t = D(c_t)h_{t-1} + D(\mathbf{W}_x x_t)(1 - c_t)$$

and

$$\begin{aligned} \frac{\partial h_t}{\partial h_{t-1}} &= D(c_t) + D(h_{t-1}) \frac{\partial c_t}{\partial h_{t-1}} + D(\mathbf{W}_x x_t) \frac{\partial(1 - c_t)}{\partial h_{t-1}} \\ &= D(c_t) + D(h_{t-1}) \frac{\partial c_t}{\partial h_{t-1}} - D(\mathbf{W}_x x_t) \frac{\partial c_t}{\partial h_{t-1}} \end{aligned}$$

and with $\sigma'(x) = \sigma(x)(1 - \sigma(x))$, we have:

$$\begin{aligned} \frac{\partial c_t}{\partial h_{t-1}} &= \frac{\partial \sigma(\mathbf{W}_c x_t + \mathbf{W}_h h_{t-1})}{\partial h_{t-1}} \\ &= D(\sigma(\mathbf{W}_c x_t + \mathbf{W}_h h_{t-1}) \odot (1 - \sigma(\mathbf{W}_c x_t + \mathbf{W}_h h_{t-1}))) \mathbf{W}_h^\top \end{aligned}$$

(d) Since ℓ is a scalar and $\mathbf{W}_x \in \mathbb{R}^{m \times n}$, $\frac{\partial \ell}{\partial \mathbf{W}_x} \in \mathbb{R}^{1 \times m \times n}$ (by numerator layout), and

$$\frac{\partial \ell}{\partial \mathbf{W}_x} \in \mathbb{R}^{1 \times m \times n} = \sum_{t=1}^K \frac{\partial \ell}{\partial h_t} \left(\frac{\partial \tilde{h}_t}{\partial \mathbf{W}_x} + \sum_{i=1}^{t-1} \left(\prod_{j=1}^{t-1} \frac{\partial h_{i+1}}{\partial h_i} \right) \frac{\partial \tilde{h}_i}{\partial \mathbf{W}_x} \right),$$

where


$$\frac{\partial \tilde{h}_t}{\partial \mathbf{W}_x} = (1 - c_t) \odot \tilde{\mathbf{W}},$$

and $\tilde{\mathbf{W}} \in \mathbb{R}^{m \times n \times m}$ such that $\tilde{\mathbf{W}}_i = [0 \dots 0 \quad x_t \quad 0 \dots 0]$ with the first zeros are of length i .

- (e)
- The network cannot have exploding gradient since the state vector does not get multiplied with a $c > 1$ constant during each time step.
 - The network can have a vanishing gradient, since σ sigmoid function can produce a value close to 0, therefore c_t is a vector of values between 0, 1. This c_t then gets Hadamard multiplied during the time step, therefore the gradient can get arbitrarily small.

1.2.2 Part 2

- (a) (Note: the recurrent gated unit graph is the same as the graph above, given the intermediate state z_t)



- (b) Replace 2 as k in the original system of equations, we have

$$\text{AttentionRNN}(k) = \begin{cases} q_0[t], q_1[t], \dots, q_k[t] &= Q_0x[t], Q_1h[t-1], \dots, Q_kh[t-k] \\ k_0[t], k_1[t], \dots, k_k[t] &= K_0x[t], K_1h[t-1], \dots, K_kh[t-k] \\ v_0[t], v_1[t], \dots, v_k[t] &= V_0x[t], V_1h[t-1], \dots, V_kh[t-k] \\ w_i[t] &= q_i[t]^\top k_i[t] \\ a[t] &= \text{softargmax}(\{w_0[t], \dots, w_k[t]\}) \\ h[t] &= \sum_{i=1}^k a_i[t]v_i[t] \end{cases}$$

- (c) In a similar fashion, $\text{AttentionRNN}(\infty)$ can be defined as $\text{AttentionRNN}(t-1)$, recurrent up to $t-1$ previous time steps at time t :

$$\text{AttentionRNN}(\infty) = \begin{cases} q_0[t], q_1[t], \dots, q_{t-1}[t] &= Q_0x[t], Q_1h[t-1], \dots, Q_{t-1}h[1] \\ k_0[t], k_1[t], \dots, k_{t-1}[t] &= K_0x[t], K_1h[t-1], \dots, K_{t-1}h[1] \\ v_0[t], v_1[t], \dots, v_{t-1}[t] &= V_0x[t], V_1h[t-1], \dots, V_kh[1] \\ w_i[t] &= q_i[t]^\top k_i[t] \\ a[t] &= \text{softargmax}(\{w_0[t], \dots, w_{t-1}[t]\}) \\ h[t] &= \sum_{i=1}^{t-1} a_i[t]v_i[t] \end{cases}$$

- (d) Since this NN has different indexing scheme compare to regular CNN, I'll use subscript notation $h[t]$ instead of underscore for h_t in part (d) of CNN question above. For $k=2$ in $\text{AttentionRNN}(2)$, we have

$$\frac{\partial h[t]}{\partial h[t-1]} = \sum_{i=0}^2 \frac{\partial(a_i[t]v_i[t])}{\partial h[t-1]},$$

where $a_i[t] = \frac{\exp(q_i[t]^\top k_0[t])}{\sum_{j=0}^2 \exp(q_j[t]^\top k_0[t])}$ by the definition of softargmax. Both $a_i[t]$ and $v_i[t]$ depends on $h[t-1]$, so we can write the above expression as

$$\frac{\partial h[t]}{\partial h[t-1]} = \sum_{i=0}^2 \frac{\partial(a_i[t]v_i[t])}{\partial h[t-1]} = \sum_{i=0}^2 \left(a_i[t] \frac{\partial v_i[t]}{\partial h[t-1]} + v_i[t] \frac{\partial a_i[t]}{\partial h[t-1]} \right)$$

Each $i = 0, 1, 2$ has different results due to dependency of a_i with respect to $h[t-1]$. By calculation, with $Z = \sum_{j=0}^2 \exp(q_j[t]^\top k_0[t])$, we have

•

$$\frac{\partial(a_0[t]v_0[t])}{\partial h[t-1]} = -v_0[t] \frac{\partial Z}{Z^2} = -v_0[t](k_1Q_q + q_1^\top K_1) \frac{\exp(q_1[t]^\top k_1[t])}{Z^2}$$

• Similarly,

$$\frac{\partial(a_2[t]v_2[t])}{\partial h[t-1]} = -v_2[t](k_1Q_q + q_1^\top K_1) \frac{\exp(q_1[t]^\top k_1[t])}{Z^2}$$

- Finally, with index 1 we can use product rule as stated above:

$$\frac{\partial(a_1[t]v_1[t])}{\partial h[t-1]} = \frac{k_1Q_1 + q_t^\top K_q \exp(q_1[t]^\top k_1[t])}{Z} - \exp(q_1[t]^\top k_1[t]) (k_1Q_q + q_1^\top K_1) \frac{\exp(q_1[t]^\top k_1[t])}{Z^2}$$

- (e) For AttentionRNN(k) during loss backpass:

$$\frac{\partial \ell}{\partial h[T]} = \sum_{t>T} \frac{\partial h[t]}{\partial h[T]} \frac{\partial \ell}{\partial h[t]}.$$

1.3 Debugging loss curve

- (a) Spike on the left due to exploding gradients.
- (b) Exploding gradients can create unstable optimization environment, e.g., let the gradient step onto the different local minima, making the loss spiking point larger than the initial loss.
- (c) We can use gradient clipping or regularization to avoid the spikes.
- (d) The model is randomly initialized so the accuracy is roughly 1/4 initially.